

Cloud **Cost** Optimization

#whoami: yogesh sharma

AWS
community
builder

- Cloud Architect @ Tata Consultancy Services
- Overall 10+ years of experience, 5+ years of experience in Cloud, DevOps
- ML and Data Enthusiast
- AWS Community Builder Program (First Batch)
- 3x AWS Certifications
- Technical Speaker and Blogger
- Owner of <https://multicloudwarrior.com/>



<https://www.linkedin.com/in/sharma-yogesh/>

What to expect

- Cloud Cost Optimization- What, Why
- Pillars for Cost Optimization
- Cost Optimization Process
- Checklist
- Latest Trends
- Q&A

We all love **Cloud and its Services...**

AWS is for builders

You build it,


You run it,

You **optimize** it

But are we fully aware about Cloud Financials

**Are we spending too much on unnecessary
stuff knowingly/unknowingly?**

Pay for what you use



Bill Formula:

$$\text{Spend} = \text{Usage} * \text{Rate}$$


But how would you ensure that you pay what you actually need?

Cloud Cost Optimization

Managing costs is a challenge for organizations using public cloud services but also an opportunity to drive efficient consumption of IT.


Top Cloud Initiatives for 2020

% of all respondents



N=750

Source: Flexera 2020 State of the Cloud Report



Pillars for Cost Optimization

Pillars for Cost Optimization



Right-sizing



Elasticity



Pricing Models



Storage Classes



Measure and Mechanisms



Designing for Cost

Right-sizing

m4.4xlarge
\$1.72 per hour


1. Provision



2. Check Metrics
(CPU, RAM, Network)

m4.large
\$0.215 per hour

3. Right-Size



4. Review Performance

86%
Saving

5. Save

Start

Choose instance that
meets your basic
requirements best

Match memory & virtual
cores

Tune

Change instance size up
or down based upon
monitoring

Use trusted advisor to
assess

Spread

Run instances across
multiple availability
zones

Smaller sizes equals
greater granularity

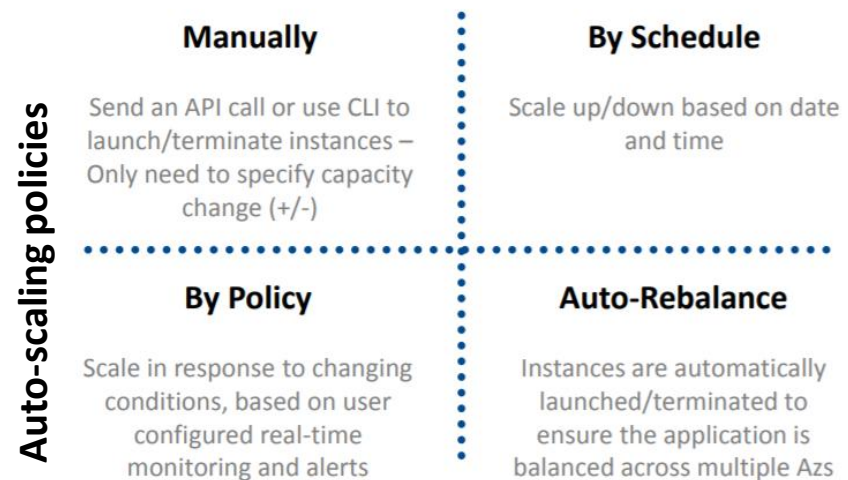
Elasticity



Think about:

- AWS Auto Scaling and EC2 Fleet
- AWS Instance Scheduler
- AWS CloudWatch


- Scale up and down to meet capacity requirements based upon CPU, RAM, network etc.
- Automatically turn nonproduction off outside of working hours



Pricing Models

On- Demand


Pay-for-compute capacity by the second with no long-term commitments



Spiky workloads to Define needs

Savings Plans & Reserved Instances

Make a commitment and receive a significant discount off compute



Committed & steady-state usage


Spot Instances

Spare AWS EC2 Capacity at saving up to 90% off On-Demand prices

If your bid > spot price → get an instance
If your bid < spot price → instance is terminated



Fault-tolerant, flexible, stateless workloads



Storage Classes

HOT




Amazon S3

WARM



Infrequent Access

COLD



Amazon Glacier

\$0.125/GB-month

Provisioned IOPS
SSD (io1)

\$0.10/GB-month
(20% cheaper)

General Purpose
SSD (gp2)


\$0.045/GB-month
(64% cheaper)

Throughput Optimized
HDD (st1)

\$0.025/GB-month
(80% cheaper)


Cold HDD (io1)

Measure and Mechanisms



**Analyze spend by
BU, account,
service, etc.**

What has changed
and why?



**Instance utilization
metrics.**

Are the instances of
the right size and
well utilized?



**Peak vs. off-peak
cost and usage.**

Are instances
turned off when they
are not needed?



**RI coverage and
utilization, Spot
usage.**

Are RIs and Spot
being used?

Designing for Cost



Consolidated Billing
Using AWS Organization



Go Serverless



Unused Elastic IPs



CloudFront




Containerization



Managed RDS

Cost Optimization Process

Cost Optimization Process




Biggest Strength- Architecture **Flexibility**

Main goal- Avoid **unnecessary** stuff



Unnecessary
Resources



Idling
Resources



Repetitive
Work

Measure: Monitor your workloads

Set up metrics to defines success and track progress

“What KPI makes sense for this workload?”




% Instances turned off daily
%instances right sized
%RI Utilization



Monitor AWS Resources




Set Alarms



Monitor Custom Metrics

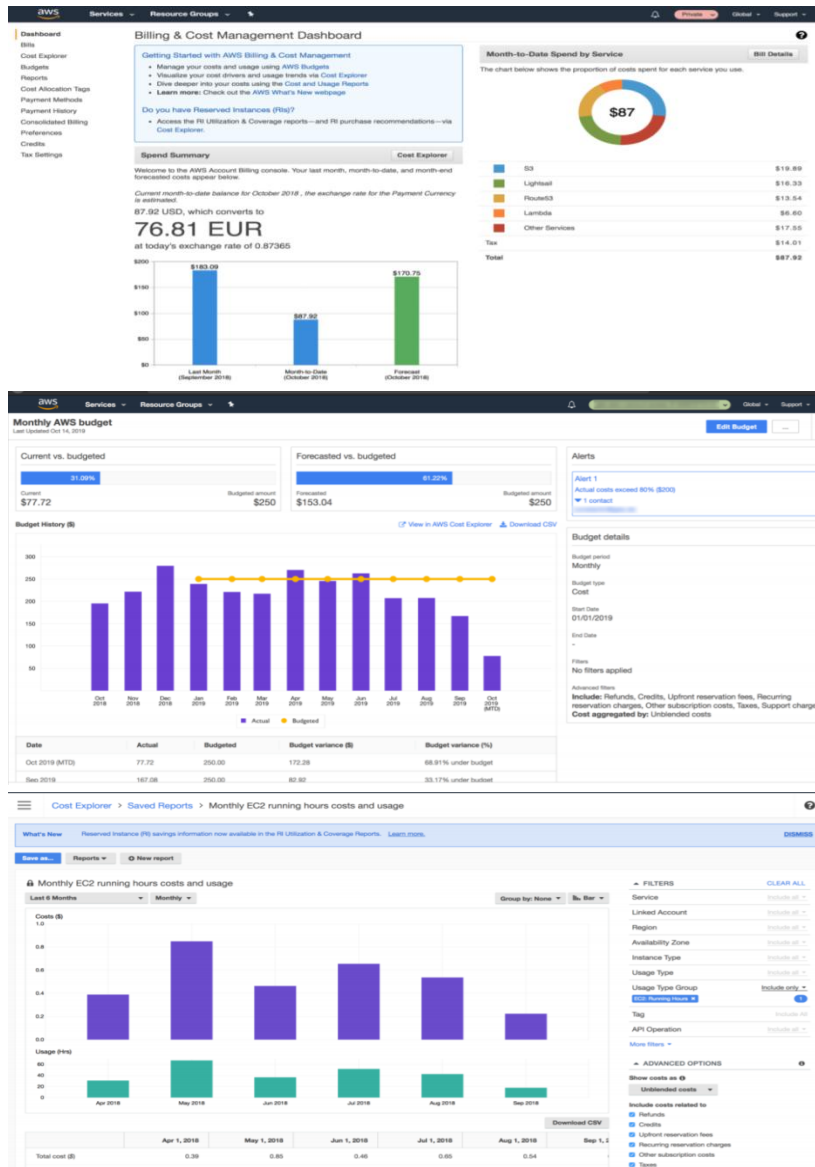


View Graphs and Statistics



Monitor & React to
Resource Changes

Measure- AWS Billing/Cost Dashboard



- Use AWS Billing and Cost Management Dashboard
- Set up your monthly AWS budget
- Dive deep with AWS Cost Explorer


Also checkout-

<https://calculator.s3.amazonaws.com/index.html>

<https://calculator.aws/>

<https://aws.amazon.com/economics/>

Measure: Cost Explorer



Use Cost Explorer Personalized Recommendations and decide to go with Reserved Instances Category-Standard, Convertible and Scheduled RIs

AWS Cost Explorer > Reserved Instance Recommendations

\$368,221

Estimated Annual Savings*

49%

Savings vs. On-Demand

39

Purchase Recommendations

Based on your past 60 days of EC2 usage, we've identified 39 three-year, all-upfront, standard RI purchase recommendations to save an estimated \$368,221 annually, representing a savings of 49% versus on-demand costs. You can take action on these recommendations in the [EC2 Reservation Purchase Console](#).

Generate recommendations based on:

All accounts Individual accounts

Sort by:

Monthly Estimated Savings

Download CSV

Purchase Recommendations (39)

Details

Buy 265 m3.medium reserved instances Size flexible**

EU (Ireland) | Linux/UNIX | Shared

Based on your past 60 days of on-demand usage, we recommend purchasing 265 m3.medium reserved instances to cover 530 normalized units per hour of m3 family usage to maximize savings.

[View Associated EC2 Usage](#)

Buy 112 m4.large reserved instances Size flexible**

EU (Ireland) | Linux/UNIX | Shared

Based on your past 60 days of on-demand usage, we recommend purchasing 112 m4.large reserved instances to cover 448 normalized units per hour of m4 family usage to maximize savings.

[View Associated EC2 Usage](#)

Monthly Estimated Savings

Upfront RI Cost

Purchase Recommendation

Instance Type

Recurring Monthly Cost: \$0.00

Expected RI Utilization: 99%

\$7,456.07

Upfront Cost:

Expected RI Utilization: 99%

Select a service

Elastic Compute Cloud (EC2)

RI Recommendation Parameters ⓘ

RI term

☐ 1 year

☒ 3 years

Offering Class

☒ Standard

☐ Convertible

Payment option

☒ All upfront

☐ Partial upfront

☐ No upfront

Based on the past

☐ 7 days

☐ 30 days

☒ 60 days

Additional Filters

Linked Account [include all](#)

AWS
community
builder

Architecture- Using AWS Organizations Consolidated Billing



Control AWS service use across accounts



Automate account creation



Consolidate billing and usage reporting

Paying Account

Linked Accounts

AWS Account 1

AWS Account 2

AWS Account 3

AWS Account 4

AWS Account 5

- **One Bill** for multiple accounts
- **Easy Tracking** of account charges (e.g., download CSV of cost data)
- **Volume Discounts** can be reached faster with combined usage
- **Reserved Instances** are shared across accounts (including RDS Reserved DBs)

Cost allocation tags

Environment

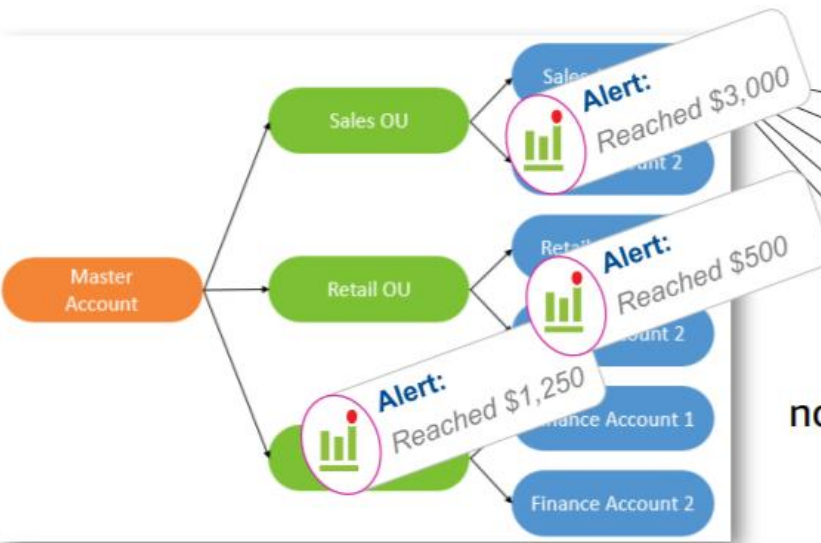
Project

Team

Application ID

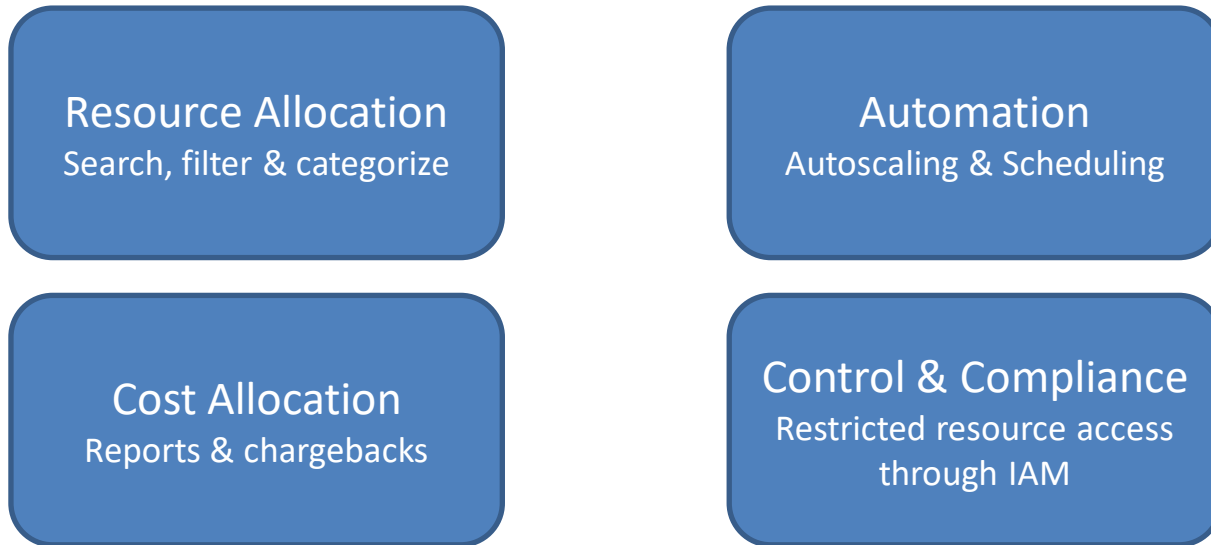
Cost center

Alerts can provide notifications when budget is reached



Architecture- Establish Tagging Strategy

What's in a Tag



Using tags to establish visibility & chargebacks

Stakeholders	Examples
<ul style="list-style-type: none">• Finance• Engineering• Line of business owners• IT• Security	<ul style="list-style-type: none">• Cost center• Application or workload• User• Expiration date• Automation support

Checklist


- Turn off unused instances- Use instance start and stop
- Delete unused/orphan volumes
- Stop paying for Idle EC2 and RDS instances and Redshift Clusters
- Understand and forecast your costs using AWS Cost Explorer and AWS Budget
- Use consolidated billing to avail volume discount across accounts
- Turn on/off whole architecture setups Using AWS CloudFormation, Terraform
- Automate Everything- Automate your instance start/stop cycle
- Choose EC2 Spot Instances
- Use Compute Savings Plan
- Use Reserved Instances
- Embrace latest trends and learn from Industry
 - Adopt Micro-service architecture
 - Modernize applications to make use of provider managed services when these are more cost-effective.
 - Consider Serverless architecture- no idle capacity
 - Adopt Containerization
 - Containers with EC2-Spot, Fargate with Spot
 - Enable S3-Intelligent Tiering- Apply Life cycle transitioning on S3 Objects

Latest Trends- Launched in re:Invent 2020

- New for AWS Lambda – 1ms Billing Granularity Adds Cost Savings
<https://aws.amazon.com/blogs/aws/new-for-aws-lambda-1ms-billing-granularity-adds-cost-savings/>
- New – Amazon EBS gp3 Volume Lets You Provision Performance Apart From Capacity
<https://aws.amazon.com/blogs/aws/new-amazon-ebs-gp3-volume-lets-you-provision-performance-separate-from-capacity-and-offers-20-lower-price/>

It lets customers independently increase IOPS and throughput without having to provision additional block storage capacity, paying only for the resources they need.

*There is no secret ingredient.
It's just you.*



Thank you!